



Red-Black Trees

Red-Black Tree

Wholeness of the Lesson

- Red-black trees provide a solution to the problem of unacceptably slow worst case performance of binary search trees. This is accomplished by introducing a new element: nodes of the tree are colored red or black, adhering to the balance condition for red-black trees. The balance condition is maintained

Red-Black Tree

during insertions and deletions and doing so introduces only slight overhead.

- **Science of Consciousness: Red-black trees, as an example of BSTs** with a balance condition, exhibit the Principle of the Second Element for solving the problem of skewed BSTs.

Review: Binary Search Trees

- *Binary Search Trees* (BSTs) are an important data structure for dynamic sets
- In addition to satellite data, elements have:
 - *key*: an identifying field inducing a total ordering
 - *left*: pointer to a left child (may be NULL)
 - *right*: pointer to a right child (may be NULL)
 - *p*: pointer to a parent node (NULL for root)


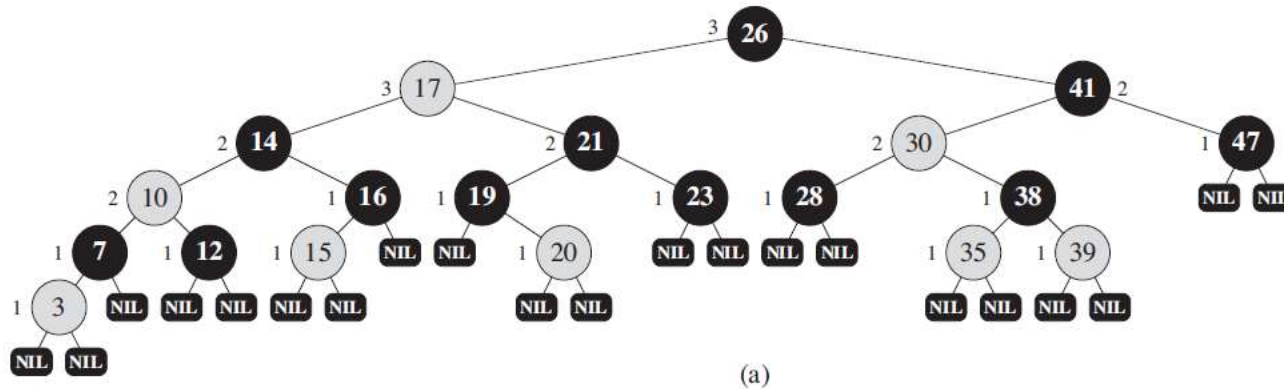
Red-Black Trees

- *Red-black trees:*
 - Binary search trees augmented with node color
 - Operations designed to guarantee that the height $h = O(\lg n)$
- First: describe the properties of red-black trees
- Then: prove that these guarantee $h = O(\lg n)$
- Finally: describe operations on red-black trees

Red-Black Properties

- The *red-black properties*:
 1. Every node is either red or black
 2. Every leaf (NULL pointer) is black
 - Note: this means every “real” node has 2 children
 3. If a node is red, both children are black
 - Note: can’t have 2 consecutive reds on a path
 4. Every path from node to descendent leaf contains the same number of black nodes
 5. The root is always black

Red-Black Tree Example



Red-Black Tree Example




Figure 13.1 A red-black tree with black nodes darkened and red nodes shaded. Every node in a red-black tree is either red or black, the children of a red node are both black, and every simple path from a node to a descendant leaf contains the same number of black nodes. (a) Every leaf, shown as a NIL, is black. Each non-NIL node is marked with its black-height; NILs have black-height 0. (b) The same red-black tree but with each NIL replaced by the single sentinel $T.nil$, which is always black, and with black-heights omitted. The root's parent is also the sentinel. (c) The same red-black tree but with leaves and the root's parent omitted entirely. We shall use this drawing style in the remainder of this chapter.

Height of Red-Black Trees

- *What is the minimum black-height of a node with height h ?*
- A: a height- h node has black-height $\geq h/2$
- Theorem: A red-black tree with n internal nodes has height $h \leq 2 \lg(n + 1)$
- *How do you suppose we'll prove this?*

Red-Black Properties: Some Proof

A red-black tree with n internal nodes has height at most $2\lg(n + 1)$.

Proof We start by showing that the subtree rooted at any node x contains at least $2^{\text{bh}(x)} - 1$ internal nodes. We prove this claim by induction on the height of x . If the height of x is 0, then x must be a leaf ($T.\text{nil}$), and the subtree rooted at x indeed contains at least $2^{\text{bh}(x)} - 1 = 2^0 - 1 = 0$ internal nodes. For the inductive step, consider a node x that has positive height and is an internal node with two children. Each child has a black-height of either $\text{bh}(x)$ or $\text{bh}(x) - 1$, depending on whether its color is red or black, respectively. Since the height of a child of x is less than the height of x itself, we can apply the inductive hypothesis to conclude that each child has at least $2^{\text{bh}(x)-1} - 1$ internal nodes. Thus, the subtree rooted at x contains at least $(2^{\text{bh}(x)-1} - 1) + (2^{\text{bh}(x)-1} - 1) + 1 = 2^{\text{bh}(x)} - 1$ internal nodes, which proves the claim.

Red-Black Properties

- The *Black Height* must be at least $h/2$. Hence,

$$n \geq 2^{h/2} - 1.$$

Moving the 1 to the left-hand side and taking logarithms on both sides yields $\lg(n + 1) \geq h/2$, or $h \leq 2\lg(n + 1)$. ■

Red-Black Tree: Some Key Cases (Rotation – for balance)



More on Rotation

LEFT-ROTATE(T, x)

```
1   $y = x.right$            // set y
2   $x.right = y.left$        // turn y's left subtree into x's right subtree
3  if  $y.left \neq T.nil$ 
4       $y.left.p = x$ 
5   $y.p = x.p$              // link x's parent to y
6  if  $x.p == T.nil$ 
7       $T.root = y$ 
8  elseif  $x == x.p.left$ 
9       $x.p.left = y$ 
10 else  $x.p.right = y$ 
11  $y.left = x$            // put x on y's left
12  $x.p = y$ 
```

More on Rotation




Figure 13.3 An example of how the procedure $\text{LEFT-ROTATE}(T, x)$ modifies a binary search tree. Inorder tree walks of the input tree and the modified tree produce the same listing of key values.

RB Insert

```
RB-INSERT(T, z)
1  y = T.nil
2  x = T.root
3  while x ≠ T.nil
4      y = x
5      if z.key < x.key
6          x = x.left
7      else x = x.right
8  z.p = y
9  if y == T.nil
10     T.root = z
11  elseif z.key < y.key
12     y.left = z
13  else y.right = z
14  z.left = T.nil
15  z.right = T.nil
16  z.color = RED
17  RB-INSERT-FIXUP(T, z)
```



RB Insert

RB-INSERT-FIXUP(T, z)


```
1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$                 // case 1
6               $y.color = BLACK$                 // case 1
7               $z.p.p.color = RED$                 // case 1
8               $z = z.p.p$                         // case 1
9          else if  $z == z.p.right$ 
10              $z = z.p$                             // case 2
11             LEFT-ROTATE( $T, z$ )                  // case 2
12              $z.p.color = BLACK$                 // case 3
13              $z.p.p.color = RED$                 // case 3
14             RIGHT-ROTATE( $T, z.p.p$ )             // case 3
15         else (same as then clause
              with “right” and “left” exchanged)
16   $T.root.color = BLACK$ 
```


RB Insert

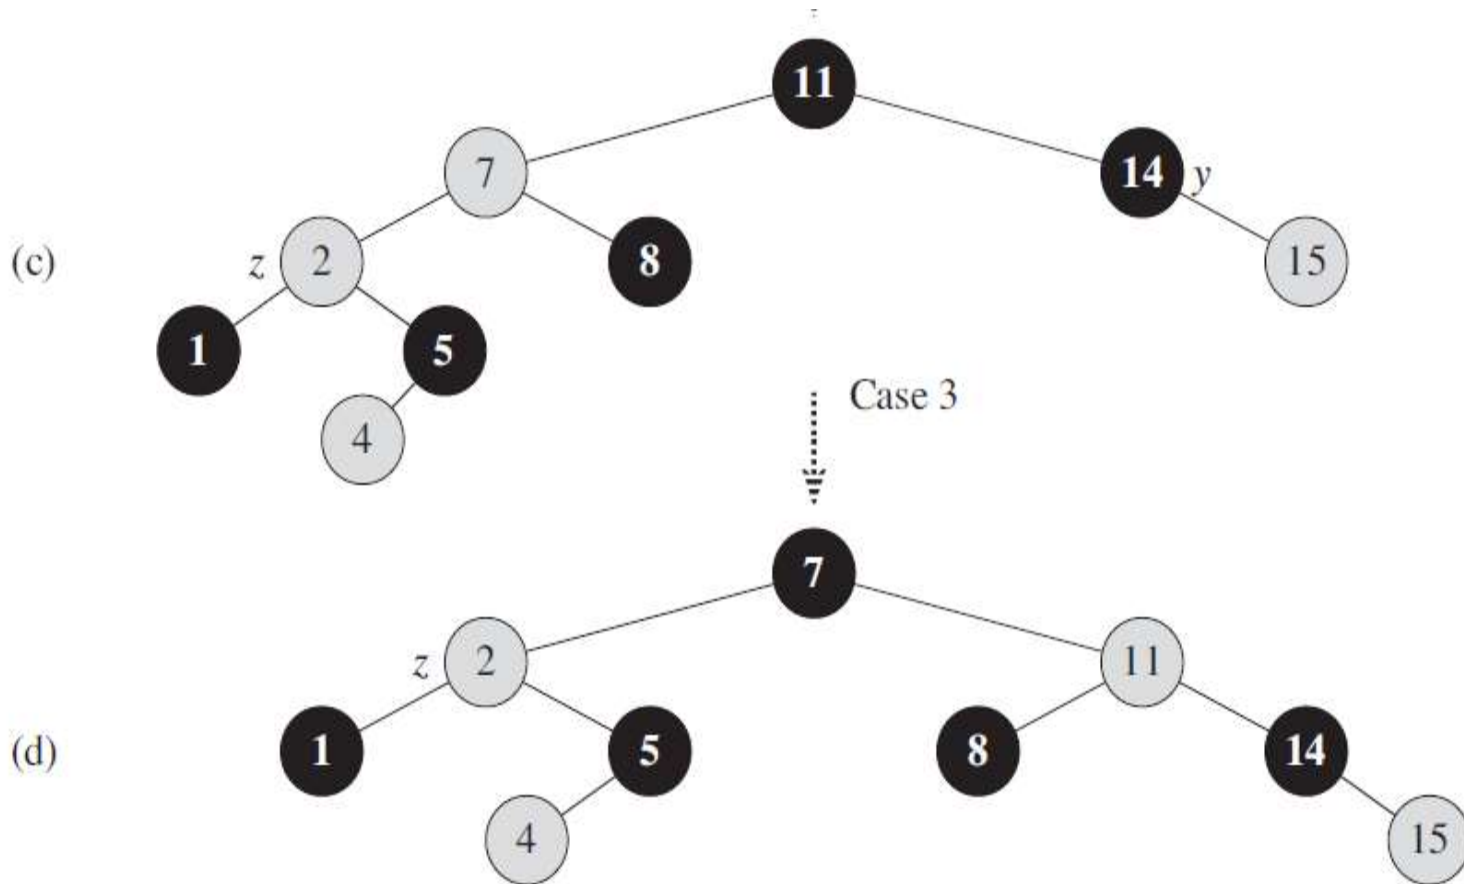
(a)



(b)



RB Insert



RB Insert




Figure 13.4 The operation of RB-INSERT-FIXUP. (a) A node z after insertion. Because both z and its parent $z.p$ are red, a violation of property 4 occurs. Since z 's uncle y is red, case 1 in the code applies. We recolor nodes and move the pointer z up the tree, resulting in the tree shown in (b). Once again, z and its parent are both red, but z 's uncle y is black. Since z is the right child of $z.p$, case 2 applies. We perform a left rotation, and the tree that results is shown in (c). Now, z is the left child of its parent, and case 3 applies. Recoloring and right rotation yield the tree in (d), which is a legal red-black tree.


RB Insert (Repeated for Conveni.)

RB-INSERT-FIXUP(T, z)

```
1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$                 // case 1
6               $y.color = BLACK$                 // case 1
7               $z.p.p.color = RED$                 // case 1
8               $z = z.p.p$                         // case 1
9          else if  $z == z.p.right$ 
10              $z = z.p$                             // case 2
11             LEFT-ROTATE( $T, z$ )                  // case 2
12              $z.p.color = BLACK$                 // case 3
13              $z.p.p.color = RED$                 // case 3
14             RIGHT-ROTATE( $T, z.p.p$ )            // case 3
15         else (same as then clause
              with “right” and “left” exchanged)
16      $T.root.color = BLACK$ 
```


Red-Black Insertion

Code – No Case



Red-Black Insertion

Code – Case 3



[Right Rotate]




Red-Black Insertion

Code – Case 3


-> Middle node
becomes root &
black

-> Root node
becomes right
node & red



Red-Black Insertion

Code – Case 1
(Just color change)




Red-Black Insertion

Code – Case 1

(Just color change

Including

`T.root.color`)



Red-Black Trees: Key Comments

- *So, the trick is to use Red and Black labels to nodes and use appropriate rules to manipulate nodes in such a way that makes the tree balanced after inserting a new value.*
- *The same is true for Deletion (we will cover it in lab problems).*

RB Trees: Worst-Case Time

- We've proved that a red-black tree has $O(\lg n)$ height
- Corollary: These operations take $O(\lg n)$ time:
 - Minimum(), Maximum()
 - Successor(), Predecessor()
 - Search()
- Insert() and Delete():
 - Will also take $O(\lg n)$ time
 - But will need special care since they modify tree